

Final Project. Geometric multigrid solver using a hierarchical mesh

Mathias Schmidt

CU Boulder's Center for Aerospace Structures develops its own Finite Element Optimization code, "Moris". This code should be able to perform Multiphysics Finite Element Analysis. One existing functionality of this code is it, to create a hierarchical mesh. The degrees of freedom are expressed by B-Spline basis. The benefit of this hierarchical structure is that the mesh can be refined according to an arbitrary criterion.

Theory hierarchical refined B-spline mesh

The finite element method is often based on Lagrange based meshes. Therefore, a given field is discretized by Lagrange interpolation functions. The degrees of freedoms (dof) are located at the nodes of the element. In areas with high field gradients a higher resolution is advantageous and thus, a refined mesh can be beneficial. B-Spline meshes are an alternative to Lagrange based meshes. B-spline meshes often need less degrees of freedom for the same accuracy than Lagrange meshes. Furthermore, they can create smooth spatial field derivatives of order $p - 1$, where p is the chosen polynomial degree of interpolation of a field. Therefore, CU Boulder's Center for Aerospace Structures developed an adaptively refining hierarchical B-spline mesh generator.

A field $\phi(x)$ is discretized with respect to a given mesh. The interpolation is given through

$$\phi(x) \approx N_i(x) \widetilde{\phi}_i \tag{0.1}$$

where $N_i(x)$ are piecewise defined polynomials, in this case B-splines, and $\widetilde{\phi}_i$ are control points that determine the weight. B-splines of the order p can be represented by $p + 2$ bases of span per spatial dimension. These smaller b-splines shall be called children. For a one dimensional spline one obtains

$$N^n = \sum_{k=0}^{p+1} w_k N_k^{n+1} \quad (0.2)$$

where n stands for the refinement level. The weights w_k can be determined by

$$w_k = \frac{1}{2^p} \binom{p+1}{k} \quad (0.3)$$

Figure 0.1. Shows such a refinement where a B-spline is represented through four weighted children B-splines. The presented example is constructed with quadratic B-splines. Higher-dimensional children weights can be created through a dyadic product of the one-dimensional ones.




Figure 0.1: Quadratic B-spline represented through sum of smaller B-splines

Figure 0.2. shows a domain discretized with B-splines. Two refinement steps are done, replacing coarse B-splines by their children representations.

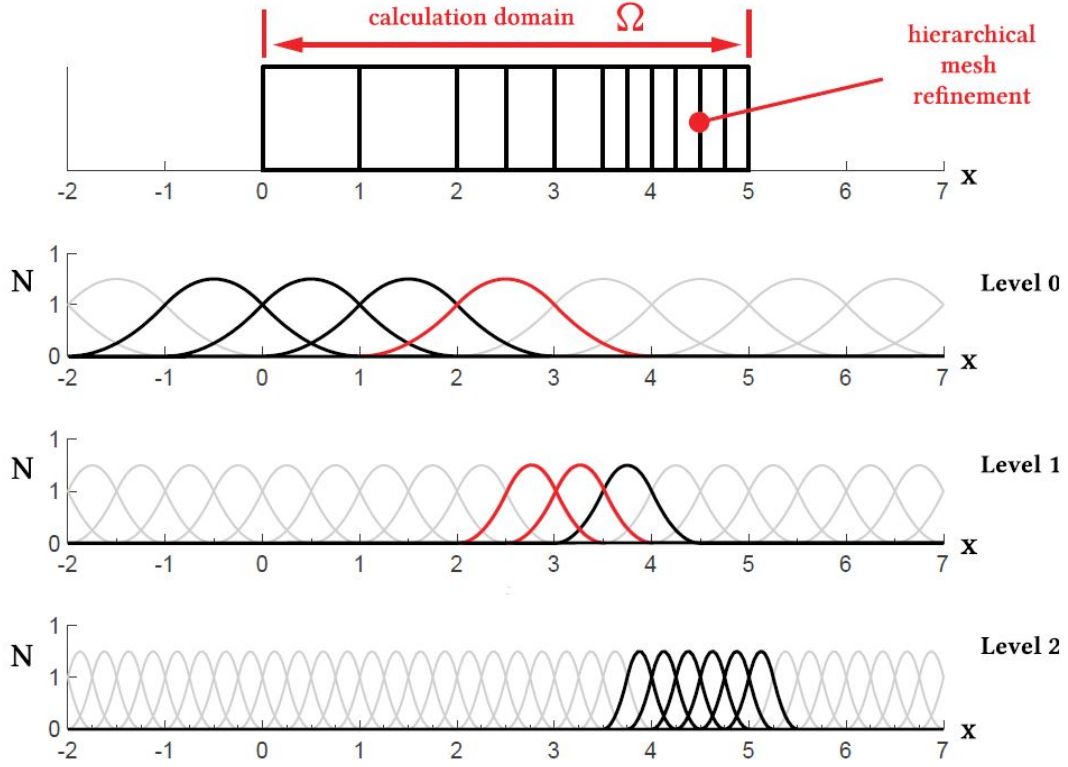


Figure 0.2: Hirarchical refined B-spline mesh

Mapping between B-spline and Lagrange mesh

A field can be discretized on a nodal basis on the mesh. A possible example for a nodal field is a signed distance field. For testing purpose, such a signed distance field will be used later on for verification of the implementation. Using a least squares projection, field data on the B-spline mesh can be mapped to a Lagrange mesh. When the B-Spline mesh is used to solve a finite element problem, the concept of element extraction plays an important role. The basic idea of element extraction is to project the B-spline degrees of freedom into a Lagrange based space, so that they can be interpreted by a Lagrange element. The Lagrange interpolation requires the field values $\hat{\phi}_j$ at the supporting Lagrange nodes \hat{x}_j . A transformation matrix T can be developed by evaluating the B-spline interpolation functions N_i at the Lagrange nodes.

$$\hat{\phi}_j = N_l(\hat{x}_j)\widetilde{\phi}_l = T_{lj}\widetilde{\phi}_l \quad (0.4)$$

Multigrid Summary

The basic idea of a multigrid method is to capture errors by utilizing multiple resolutions. Oscillatory components are effectively reduced through a simple iterative procedure, while




Figure 0.3: Multigrid V-cycle

smooth components are tackled using auxiliary lower resolution versions of the problem. For a multigrid preconditioner a hierarchy of grids is constructed. Then, a smoother such as the Jacobi or Gauss-Seidel method is applied to the linear system associated with the finest grid. Next, using a Galerkin approach, residual and Jacobian are restricted to the second finest grid, and a smoother is applied to the residual equation associated with this grid. This process is repeated until we arrive at the coarsest grid. On the coarsest grid a direct solver is performed. This coarse grid solution is prolonged back to the finest grid. A smoothing is performed on each level. This process is called the V-cycle. The V-cycle can be applied until converging to a solution or it can be applied as a preconditioner to another solver. The V-cycle is illustrated in Figure 0.3.

Code implementation

Hierarchical mesh refinement (HMR) implementation

The idea of this project is to exploit the hierarchical nature of the B-spline basis to create prolongation and restriction operators which can then be used to create a geometric multigrid preconditioner. Therefore, the hierarchical mesh generator must be modified in a way that it calculates the projection stencils corresponding to the pairs of fine and coarse B-spline basis. This is done in the appended file `cl_HMR_Bspline_Mesh_Base.cpp`. The function `BSpline_Mesh_Base::calculate_child_stencil()` calculates a non-truncated stencil for two and three dimensions. This stencil is adjusted for every parent basis and can be requested using `BSpline_Mesh_Base::get_children_ind_for_basis()` and `BSpline_Mesh_Base::get_children_weights_for_parent_basis()`. The first function returns the children indices. The second returns the corresponding weights.

Model Solver Interface (MSI) implementation

The Model Solver Interface serves as the interface between the finite element model and the nonlinear or time solver. This module enables building and solving the following system of nonlinear equations:

$$F_i(u_k(a_j)) = 0 \quad (0.5)$$

where F_i is the residual equation, a_j are independent abstract variables and u_k are dependent physical variables. It is assumed that the residuals F_i are functions of physical variables, i.e. physical degrees of freedom (pdofs), which in turn depend on abstract degrees of freedom (adofs). The pdofs are defined by a type, such as temperature or displacement in x-direction, and a vector of time step indices. Adofs are derived from pdofs using information provided by the transformation matrices presented above. Each adof is linked to one type and time step ID.

To create a multigrid preconditioner, the model solver interface has to provide a mapping between the different degrees of freedom on the multigrid levels, for each dof type and time. This is done in the file `cl_MSI_Multigrid.cpp`.

Linear Solver (DLA) multigrid implementation

In order to use PETSc, wrappers are written around the PETSc matrix, vector, and the linear solver. This was done to fit PETSc into the existing framework. In addition, the interpolation operators are assembled in `cl_DLA_Geometric_Multigrid.cpp`. The interpolation operators are assembled using the weights provided by the hierarchical mesh and the mapping provided by the model solver interface. The operators are provided to PETSc's PCMG module to build a multigrid preconditioner. This is done in the file `cl_DLA_Linear_Solver_PETSc.cpp`.

Evaluation and Results

For debugging purposes two simple hierarchical meshes are created. The first mesh operates on linear B-splines, the second one on quadratic B-splines. In both meshes the bottom left element is refined twice. The meshes and the corresponding B-spline basis are presented in Figure 0.4. The top row shows linear B-spline basis. The bottom row quadratic B-spline basis. In the figure to the right the B-spline basis on different levels are warped to show the levels. Red basis are inactive. Blue ones are active on the fines level. Green ones are inactive on the fines level but will be used when a coarser level is reached during the multigrid process. The presented mesh on the plane is the corresponding Lagrange mesh. The different refinement patterns between linear and quadratic plot are based on a refinement




Figure 0.4: Top: Linear B-spline basis. Bottom: quadratic B-spline Basis. Linear and quadratic basis are presented in a view from above and from an angle. The B-spline basis in different levels are warped to present the levels.

buffer criterion which increases with increasing B-spline order. The shown meshes have two coarser multigrid meshes.

The created prolongation or restriction operators for these B-Spline meshes can be printed to screen and compared to the operator created by hand. This was done and the following two operators were found for the linear case. The found operators correspond to the expected result. The operators are presented in Figure 0.5.

The same was done for the operators based on quadratic B-splined. The sparse matrices correspond the expected results again. The operators are presented in Figure 0.6.

```

Mat Object: 1 MPI processes
  type: seqaij
row 0: (4, 1.)
row 1: (5, 1.)
row 2: (6, 1.)
row 3: (7, 1.)
row 4: (8, 1.)
row 5: (9, 1.)
row 6: (10, 1.)
row 7: (11, 1.)
row 8: (12, 1.)
row 9: (13, 1.)
row 10: (14, 1.)
row 11: (15, 1.)
row 12: (16, 1.)
row 13: (17, 1.)
row 14: (18, 1.)
row 15: (19, 1.)
row 16: (20, 1.)
row 17: (21, 1.)
row 18: (22, 1.)
row 19: (0, 1.) (1, 0.5) (2, 0.5) (3, 0.25)

Mat Object: 1 MPI processes
  type: seqaij
row 0: (13, 1.)
row 1: (14, 1.)
row 2: (17, 1.)
row 3: (18, 1.)
row 4: (0, 0.5) (1, 0.5) (2, 0.25) (19, 1.)
row 5: (0, 0.5) (2, 0.25) (3, 1.) (4, 0.5) (5, 0.5) (6, 0.25)
row 6: (1, 0.5) (2, 0.25) (9, 1.) (10, 0.5) (11, 0.5) (12, 0.25)
row 7: (4, 0.5) (6, 0.25) (7, 1.) (8, 0.5)
row 8: (11, 0.5) (12, 0.25) (15, 1.) (16, 0.5)

```

Figure 0.5: Multigrid operators for the linear case. Left: First coarsening step. Right: Second coarsening step

```

Mat Object: 1 MPI processes
  type: seqaij
row 0: (4, 1.)
row 1: (5, 1.)
row 2: (6, 1.)
row 3: (7, 1.)
row 4: (8, 1.)
row 5: (9, 1.)
row 6: (10, 1.)
row 7: (11, 1.)
row 8: (12, 1.)
row 9: (13, 1.)
row 10: (14, 1.)
row 11: (15, 1.)
row 12: (16, 1.)
row 13: (17, 1.)
row 14: (18, 1.)
row 15: (19, 1.)
row 16: (20, 1.)
row 17: (21, 1.)
row 18: (22, 1.)
row 19: (23, 1.)
row 20: (24, 1.)
row 21: (25, 1.)
row 22: (26, 1.)
row 23: (27, 1.)
row 24: (28, 1.)
row 25: (29, 1.)
row 26: (30, 1.)
row 27: (31, 1.)
row 28: (32, 1.)
row 29: (33, 1.)
row 30: (34, 1.)
row 31: (35, 1.)
row 32: (36, 1.)
row 33: (37, 1.)
row 34: (38, 1.)
row 35: (0, 0.5625) (1, 0.1875) (2, 0.1875) (3, 0.0625)

Mat Object: 1 MPI processes
  type: seqaij
row 0: (0, 0.1875) (1, 0.1875) (2, 0.0625) (35, 0.5625)
row 1: (0, 0.5625) (1, 0.0625) (2, 0.1875) (3, 0.5625) (4, 0.1875) (5, 0.1875) (6, 0.0625) (35, 0.1875)
row 2: (0, 0.0625) (1, 0.5625) (2, 0.1875) (11, 0.5625) (12, 0.1875) (13, 0.1875) (14, 0.0625) (35, 0.1875)
row 3: (0, 0.1875) (1, 0.1875) (2, 0.5625) (3, 0.1875) (4, 0.0625) (5, 0.5625) (6, 0.1875) (11, 0.1875)
row 4: (12, 0.5625) (13, 0.0625) (14, 0.1875) (15, 0.5625) (16, 0.1875) (17, 0.1875) (18, 0.0625) (35, 0.0625)
row 5: (3, 0.1875) (4, 0.5625) (5, 0.0625) (6, 0.1875) (7, 0.5625) (8, 0.1875) (9, 0.1875) (10, 0.0625)
row 6: (3, 0.0625) (4, 0.1875) (5, 0.1875) (6, 0.5625) (7, 0.1875) (8, 0.0625) (9, 0.5625) (10, 0.1875)
row 7: (15, 0.1875) (16, 0.5625) (17, 0.0625) (18, 0.1875) (19, 0.5625) (20, 0.1875) (21, 0.1875) (22, 0.0625)
row 8: (7, 0.1875) (8, 0.5625) (9, 0.0625) (10, 0.1875)
row 9: (7, 0.0625) (8, 0.1875) (9, 0.1875) (10, 0.5625) (19, 0.1875) (20, 0.5625) (21, 0.0625) (22, 0.1875)
row 10: (11, 0.1875) (12, 0.0625) (13, 0.5625) (14, 0.1875) (23, 0.5625) (24, 0.1875) (25, 0.1875) (26, 0.0625)
row 11: (11, 0.0625) (12, 0.1875) (13, 0.1875) (14, 0.5625) (15, 0.1875) (16, 0.0625) (17, 0.5625) (18, 0.1875)
row 12: (23, 0.1875) (24, 0.5625) (25, 0.0625) (26, 0.1875) (27, 0.5625) (28, 0.1875) (29, 0.1875) (30, 0.0625)
row 13: (15, 0.0625) (16, 0.1875) (17, 0.1875) (18, 0.5625) (19, 0.1875) (20, 0.0625) (21, 0.5625) (22, 0.1875)
row 14: (27, 0.1875) (28, 0.5625) (29, 0.0625) (30, 0.1875) (31, 0.5625) (32, 0.1875) (33, 0.1875) (34, 0.0625)
row 15: (19, 0.0625) (20, 0.1875) (21, 0.1875) (22, 0.5625) (31, 0.1875) (32, 0.5625) (33, 0.0625) (34, 0.1875)
row 16: (23, 0.1875) (24, 0.0625) (25, 0.5625) (26, 0.1875)
row 17: (23, 0.0625) (24, 0.1875) (25, 0.1875) (26, 0.5625) (27, 0.1875) (28, 0.0625) (29, 0.5625) (30, 0.1875)
row 18: (27, 0.0625) (28, 0.1875) (29, 0.1875) (30, 0.5625) (31, 0.1875) (32, 0.0625) (33, 0.5625) (34, 0.1875)
row 19: (31, 0.0625) (32, 0.1875) (33, 0.1875) (34, 0.5625)

```

Figure 0.6: Multigrid operators for the quadratic case. Left: First coarsening step. Right: Second coarsening step

Multigrid Solver

The effectiveness of the multigrid preconditioner is compared to an ILU preconditioner. Therefore, the following settings for the ILU preconditioner and the multigrid preconditioner were chosen. As a solver a flexible gmres is used.

ILU-Preconditioner

- Level of fill: 0
- Drop tolerance: 10^{-6}

Multigrid Preconditioner

- Two coarser multigrid level
- Multigrid form: Multiplicative
- The coarser grids are computed from the fines grid using the Galerkin process
- A preconditioned Richardson iterative method is used to solve the coarsest level
- As a smoother down and up the V-cycle, one iteration of a GMRES is used with a jacobi preconditioner

A two-dimensional problem and a three-dimensional problem are used. The problem is set up in the test `cl_DLA_Multigrid.Test.cpp`. The experiment was repeated while the underlying tensor grid was refined for both the two- and three-dimensional problem. The signed distance field which is mapped by the least squares projection is defined as followed.

$$v = ||\vec{x}||_2 - 0.9 \quad (0.6)$$

Two-Dimensional Mapping Problem

The finest two-dimensional problem is presented in Figure 0.7. The domain has the size 1 by 1 unit length. The underlying tensor grid has 32x32 elements. After refining 19782 elements are existent. Two refienment steps are performed on a circle with a radius of 0.9. The signed distance field is evaluated on the B-spline mesh. It is mapped to the Lagrange mesh. The result is presented in Figure 0.7. Using the multigrid preconditioner delivers the same result vector than using the ILU preconditioner. The colors shown in Figure 0.7 show the values of the mapped signed distance field.

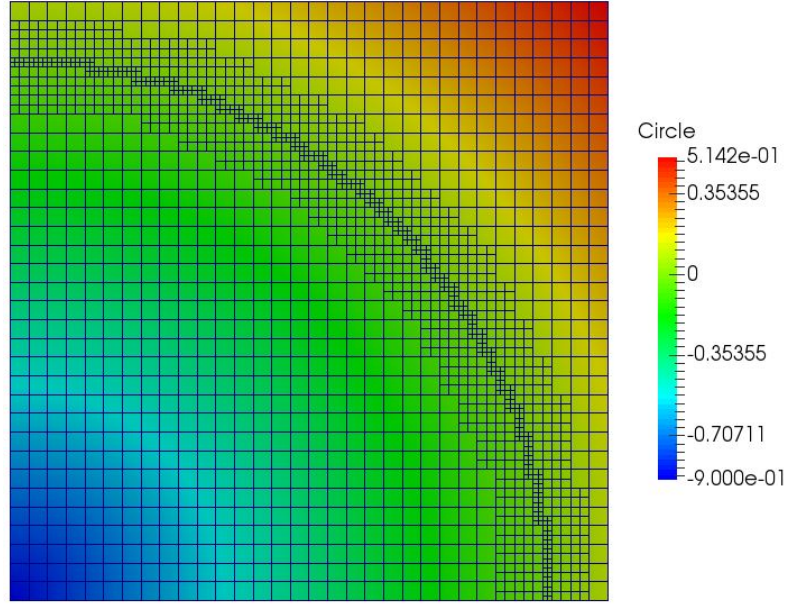


Figure 0.7: 2D-example. The colorbar shows the values of the mapped signed distance field

For the presented solver settings, the comparison of the multigrid preconditioner and the ILU shows for the solver with the multigrid preconditioner needs 3 to 4 iterations to solve the problem while the solver with the ILU preconditioner need 6 iterations. The number of iterations needed stay nearly constant while increasing the mesh size

Three-Dimensional Mapping Problem

The same analysis as for the two-dimensional problem was done for the three-dimensional problem shown in Figure 0.9. The domain has the size 1 by 1 by 1 unit length. Two refinement steps are performed on a sphere with a radius of 0.9. In Figure 0.10 a plot is presented which compares the number of iterations needed to solve the three-dimensional problem with either an ILU or a multigrid preconditioner. Using the multigrid preconditioner delivers the same result vector than using the ILU preconditioner. The colors shown in Figure 0.9 show the values of the mapped signed distance field. Results are presented for different mesh refinements. For the presented solver settings, the comparison of the multigrid preconditioner and the ILU shows for the solver with the multigrid preconditioner needs 4 to 6 iterations to solve the problem while the solver with the ILU preconditioner need 6 to 7 iterations. The number of iterations needed stay nearly constant while increasing the mesh size.




Figure 0.8: Comparison number of iterations needed for two dimensional problem

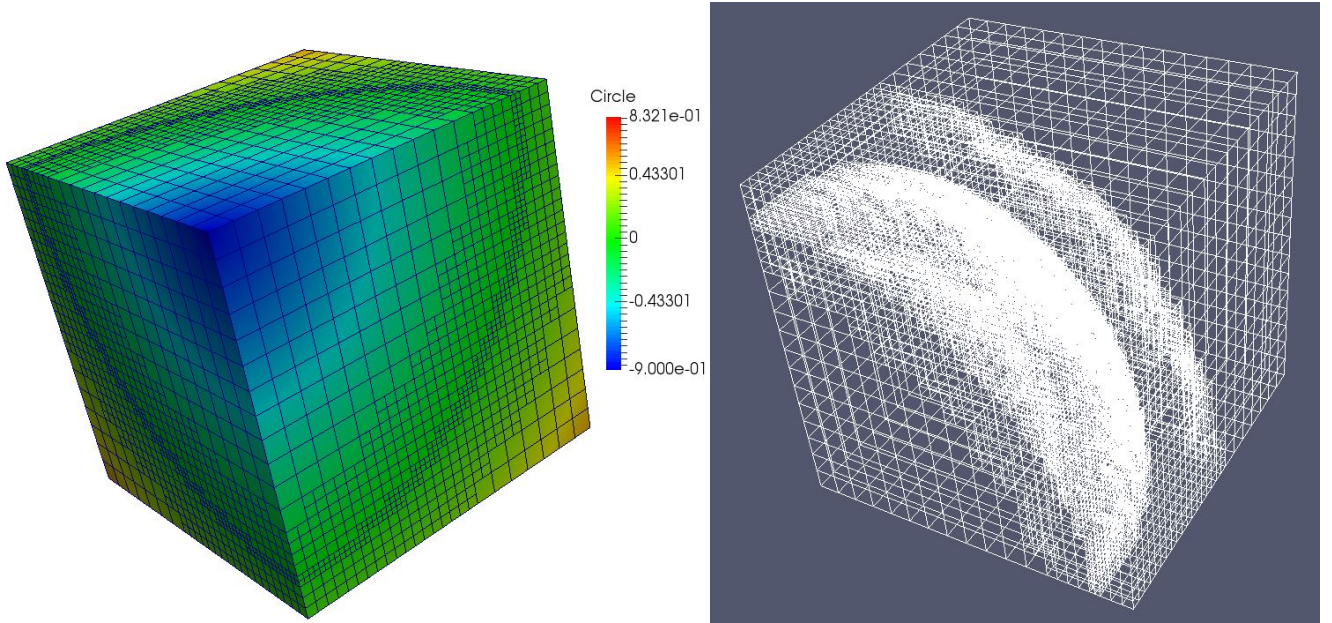


Figure 0.9: 3D-example. Left: The colorbar shows the values of the mapped signed distance field. Right: Wireframe mesh




Figure 0.10: Comparison number of iterations needed for two dimensional problem

Conclusion

It was shown that a multigrid preconditioner based on the hierarchical structure of the hierarchical B-spline mesh can be used as a preconditioner for a Krylov solver. To see the potential of this preconditioner more solver option must be studied. As for example other types of smoothers. Furthermore, the implemented stencil which maps from coarse to fine B-spline basis is not truncated. Studies for a truncated stencil should be performed to see its effect. Moreover, the code has to be profiled to find potential bottle necks. Due to the limited amount of time, this was not done yet.